Ruchira Dhar

Email | Website | LinkedIn | Google Scholar

Education

PhD in Computer Science 2023 – Present
Department of Computer Science, University of Copenhagen, Denmark Lab Website
Supervisor: Dr. Anders Søgaard · *Focus*: Science of Evaluation, Multilingual Evaluation

Visiting Researcher Jun – Jul 2026
Department of Computer Science, ETH Zurich, Switzerland Lab Website
Supervisor: Dr. Mrinmaya Sachan · *Focus*: Psychometrics-based Evaluation of Language Models

Masters in Linguistics 2018 – 2020
Department of Linguistics, English and Foreign Languages University, Hyderabad, India
Supervisor: Dr. Utpal Lahiri · *Thesis*: Formal Semantics of Adjective-Noun Composition

Experience

AI Research Engineer Intern, Refractal Aug 2026 – Present
London, UK

- Researching Privacy-Detection, Jailbreak and Harmful Task Evaluations
- Building evaluation harness and model monitoring pipeline for Agentic AI

Member of Technical Staff, Flowrite Oy Feb 2022 – Jun 2023
Helsinki, Finland

- Research in model evaluation pipelines for generative LLMs
- Developed large-scale annotation and user feedback frameworks for model alignment

NLG Engineer, Cerence Inc May 2021 – Feb 2022
Pune, India

- Engineered NLG and Speech Recognition for multilingual voice assistants
- Conducted evaluation for speech-driven interactions in real-world automotive settings

Associate Engineer, Kore AI Jan – May 2021
Hyderabad, India

- Developed production-ready chatbot pipelines leveraging intent classification and entity recognition
- Collaborated with cross-functional teams to deliver AI-powered tools across different interfaces

Research Fellow, Indian Institute of Technology (IIT) Sep – Dec 2020
Delhi, India

- Advanced research on multilingual semantic annotation pipelines in NLP

First-Author Publications

Does Item Response Theory Generalize Across Languages? Evidence from Cross-Lingual LLM Benchmarks
Ruchira Dhar, Qiwei Peng, Anders Søgaard.
Findings of ACL: Empirical Methods in Natural Language Processing (EMNLP), ACL Anthology, 2026.
Paper Code

Evaluating Adjective-Noun Compositionality in LLMs: Functional vs Representational Perspectives
Ruchira Dhar, Qiwei Peng, Anders Søgaard.
*Proceedings of The 15th Joint Conference on Lexical and Computational Semantics (*SEM), ACL Anthology*, 2026.
Paper Code

Beyond Technocratic XAI: The Who, What & How in Explanation Design
Ruchira Dhar, Stephanie Brandl, Ninell Oldenburg, Anders Søgaard.
Proceedings of The AAAI/ACM Conference on AI, Ethics, & Society (AIES), 2025.
Paper

Investigating the role of modality and training objective on representational alignment between transformers and the brain
Hyewon Willow Han*, **Ruchira Dhar***, Qingqing Yang*, et al.
Proceedings of UniReps: the Second Edition of the Workshop on Unifying Representations in Neural Models at NeurIPS, PMLR 285, 2024. (*Equal Contribution First Authors)
Paper

First-Author Preprints

On Challenges in Evaluation Methodology: A Review
Ruchira Dhar, Qiwei Peng, Anders Søgaard.
arXiv Preprint, 2026.
Paper

On the Measure of a Model: From Intelligence to Generality
Ruchira Dhar, Ninell Oldenburg, Anders Søgaard.
Presented at Science of Benchmarking and Evaluating AI Workshop, EurIPS, 2025.
Paper

EvalCards: A Framework for Standardized Evaluation Reporting
Ruchira Dhar, et al.
Presented at Science of Benchmarking and Evaluating AI Workshop, EurIPS, 2025.
Paper

Contributing Publications

CulturalRiddles: A Multicultural Benchmark of Riddles
My Chiffon Nguyen, Tianyi Hu, . . . , **Ruchira Dhar**, et al.
6th Multilingual Representation Learning (MRL) Workshop at EMNLP, Budapest, Hungary, 2026.
Paper

Evaluation Cards: An Interpretive Layer for AI Evaluation Reporting
Avijit Ghosh, Anka Reuel, . . . , **Ruchira Dhar**, et al.
Forthcoming in *Proceedings of NeurIPS, Evaluations & Datasets Track*, 2026. **Spotlight**.
Paper

Mechanistic Interpretability Needs Philosophy
Iwan Williams, Ninell Oldenburg, **Ruchira Dhar**, et al.
Forthcoming in *Proceedings of NeurIPS, Position Paper Track*, 2026. **Spotlight**.
Paper

Realist and Pluralist Conceptions of Intelligence and Their Implications on AI Research
Ninell Oldenburg, **Ruchira Dhar**, Anders Søgaard.
Proceedings of the Fortieth AAAI Conference on Artificial Intelligence, 2026. **Oral**.
Paper

Evaluating Multimodal Language Models as Visual Assistants for Visually Impaired Users
Antonia Karamolegkou, Malvina Nikandrou, Georgios Pantazopoulos, Danae Sanchez Villegas, Phillip Rust, **Ruchira Dhar**, Daniel Hershcovich, Anders Søgaard.
Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics, ACL Anthology, 2025.
Paper

Defining Knowledge: Bridging Epistemology and Large Language Models
Constanza Fierro, **Ruchira Dhar**, Filippos Stamatiou, Nicolas Garneau, Anders Søgaard.
2024 Proceedings of Empirical Methods in Natural Language Processing (EMNLP) Main, ACL Anthology, 2024.
Paper

Invited Talks

"Are We Testing the Right Things in Multi-Agent AI?". UKAI Forum & Meridian Cambridge, St John's College, University of Cambridge, United Kingdom. October 2025.
Slides

"The Completeness Problem: Beyond Human Metrics in Assessing Abilities of Cognitive Systems.". International Symposium on AI & Neuroscience (ACAIN), Italy. September 2024.
Slides

Teaching

Teaching Assistant — NLP Course (NDAK18000U) 2023 – Present
Department of Computer Science, University of Copenhagen
Undergraduate and graduate-level instruction in natural language processing.

Academic Service

Program Committee Member, Measuring What Matters, WASP, 2026.
Website

Reviewer: ACL, EMNLP, NeurIPS, ICML.

External Affiliations

Research Fellow, EvalEval Coalition 2025 – Present

Invited to contribute to research on AI evaluation and reporting, including user research and interviews informing interface design. Contributor in the *Evaluation Cards* initiative.

Website

Open Science Lead, Cohere Labs 2024 – Present

Co-lead the Research Connections programme, organising research talks and collaboration sessions for a global community of 200+ researchers.

Website

Fellow, Centre for Philosophy of AI, University of Copenhagen 2023 – Present

Lead interdisciplinary research on the ethical and societal implications of AI, connecting model evaluation and interpretability with questions of knowledge, accountability and AI governance.

Website

Skills

Programming Languages: Python, R, C++, SQL

Machine Learning and NLP: PyTorch, Hugging Face, spaCy

Data Analysis and Scientific Computing: NumPy, Pandas, Jupyter Notebook, Google Colab

Experiment Tracking and Collaboration: Weights & Biases, Git, GitHub, Jira, Confluence

Human Data Collection: Prolific

Scientific Writing: LaTeX, Markdown

Languages: Danish (Limited Working Proficiency), English (Native Proficiency), Bengali (Native Proficiency), Hindi (Native Proficiency)